

Word Embeddings

Presented By

Deepak Dhingan, Roll No. 25524007

Department of IT, NIT Jalandhar

Cornell Bowers C-IS



What is an easy to make pasta recipe that features garlic and red sauce?



Here is a simple pasta recipe that features garlic and red sauce:

Ingredients:

- * 1 pound pasta (such as spaghetti or penne)
- * 1 tablespoon olive oil
- * 2 cloves garlic, minced
- * 1 (28-ounce) can crushed tomatoes
- * 1/4 teaspoon salt
- * 1/4 teaspoon black pepper
- * 1/4 teaspoon red pepper flakes (optional)
- * 1/2 cup grated Parmesan cheese
- * Fresh parsley or basil, chopped, for garnish (optional)



san f

san francisco weather
san francisco
san francisco giants
san fernando valley
san francisco state university
san francisco hotels
san francisco 49ers
san fernando
san fernando mission
san francisco zip code

Google Search

I'm Feeling Lucky

Spanish - detected



English

¿Cómo estás?



How are you?



Verified

[Open in Google Translate](#) • [Feedback](#)

```
// write a binary search algorithm
const binarySearch = (arr, target) => {
  let left = 0;
  let right = arr.length - 1;
  let middle = Math.floor((left + right) / 2);
  while (arr[middle] !== target && left <= right) {
    if (target < arr[middle]) {
      right = middle - 1;
    } else {
      left = middle + 1;
    }
    middle = Math.floor((left + right) / 2);
  }
  return arr[middle] === target ? middle : -1;
}
```

Representing words by their context

Distributional hypothesis: words that occur in similar contexts tend to have similar meanings



J.R.Firth 1957

- “You shall know a word by the company it keeps”
- One of the most successful ideas of modern statistical NLP!

*...government debt problems turning into **banking** crises as happened in 2009...*

*...saying that Europe needs unified **banking** regulation to replace the hodgepodge...*

*...India has just given its **banking** system a shot in the arm...*

These context words will represent *banking*.

Distributional hypothesis

“tejuino”



C1: A bottle of ____ is on the table.

C2: Everybody likes ____.

C3: Don't have ____ before you drive.

C4: We make ____ out of corn.

Distributional hypothesis

C1: A bottle of ____ is on the table.

C2: Everybody likes ____.

C3: Don't have ____ before you drive.

C4: We make ____ out of corn.

	C1	C2	C3	C4
tejuino	1	1	1	1
loud	0	0	0	0
motor-oil	1	0	0	0
tortillas	0	1	0	1
choices	0	1	0	0
wine	1	1	1	0

“words that occur in similar contexts tend to have similar meanings”

How to represent words?

N-gram language models

It is 76 F and ____.

$$P(w \mid \text{it is 76 F and})$$

[0.0001, 0.1, 0, 0, 0.002, ..., 0.3, ..., 0]
red sunny

Text classification

I like this movie. 👍

$w^{(1)}$ [0, 1, 0, 0, 0, ..., 1, ..., 1]

I don't like this movie. 👎

$w^{(2)}$ [0, 1, 0, 1, 0, ..., 1, ..., 1]
don't

Representing words as discrete symbols

In traditional NLP, we regard words as discrete symbols:

hotel, conference, motel — a localist representation

one 1, the rest 0's



Words can be represented by one-hot vectors:

hotel = [0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]

motel = [0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0]

Vector dimension = number of words in vocabulary (e.g., 500,000)


Challenge: How to compute similarity of two words?

Words as vectors

- We'll build a new model of meaning focusing on similarity
 - Each word is a vector
 - Similar words are “nearby in space”
- A first solution: we can just use context vectors to represent the meaning of words!
 - word-word co-occurrence matrix:

	aardvark	computer	data	pinch	result	sugar	...
apricot	0	0	0	1	0	1	
pineapple	0	0	0	1	0	1	
digital	0	2	1	0	1	0	
information	0	1	6	0	4	0	

Words as vectors



$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$

$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\sum_{i=1}^V u_i v_i}{\sqrt{\sum_{i=1}^V u_i^2} \sqrt{\sum_{i=1}^V v_i^2}}$$

What is the range of $\cos(\cdot)$?

Words as vectors

Problem: not all counts are equal, words can randomly co-occur

- Solution: re-weight by how likely it is for the two words to co-occur by simple chance
- PPMI = Positive Pointwise Mutual Information

$$\text{PPMI}(w, c) = \max(\log_2 \frac{P(w, c)}{P(w)P(c)}, 0)$$

	computer	data	result	pie	sugar
cherry	2	8	9	442	25
strawberry	0	0	1	60	19
digital	1670	1683	85	5	4
information	3325	3982	378	5	13

	computer	data	result	pie	sugar
cherry	0	0	0	4.38	3.30
strawberry	0	0	0	4.10	5.51
digital	0.18	0.01	0	0	0
information	0.02	0.09	0.28	0	0

Sparse vs dense vectors

- Still, the vectors we get from word-word occurrence matrix are sparse (most are 0's) & long (vocabulary size)
- Alternative: we want to represent words as **short** (50-300 dimensional) & **dense** (real-valued) vectors
 - The focus of this lecture
 - The basis of all the modern NLP systems

Dense vectors

$$\text{employees} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 10.109 \\ -0.542 \\ 0.349 \\ 0.271 \\ 0.487 \end{pmatrix}$$



Why dense vectors?

- Short vectors are easier to use as features in ML systems
- Dense vectors may generalize better than storing explicit counts
- They do better at capturing synonymy
 - w_1 co-occurs with “car”, w_2 co-occurs with “automobile”
- Different methods for getting dense vectors:
 - Singular value decomposition (SVD)
 - **word2vec and friends: “learn” the vectors!**

SVD

$$\begin{array}{|c|} \hline \text{word-word} \\ \text{PPM matrix} \\ \hline X \\ \hline w \times c \\ \hline \end{array} = \begin{array}{|c|} \hline W \\ \hline w \times m \\ \hline \end{array} \begin{array}{|c|} \hline \Sigma \\ \hline m \times m \\ \hline \end{array} \begin{array}{|c|} \hline C \\ \hline m \times c \\ \hline \end{array}$$

Word2vec

- Input: a large text corpora, V, d

- V : a pre-defined vocabulary
- d : dimension of word vectors (e.g. 300)
- Text corpora:
 - Wikipedia + Gigaword 5: 6B
 - Twitter: 27B
 - Common Crawl: 840B

$$v_{\text{cat}} = \begin{pmatrix} -0.224 \\ 0.130 \\ -0.290 \\ 0.276 \end{pmatrix} \quad v_{\text{dog}} = \begin{pmatrix} -0.124 \\ 0.430 \\ -0.200 \\ 0.329 \end{pmatrix}$$

$$v_{\text{the}} = \begin{pmatrix} 0.234 \\ 0.266 \\ 0.239 \\ -0.199 \end{pmatrix} \quad v_{\text{language}} = \begin{pmatrix} 0.290 \\ -0.441 \\ 0.762 \\ 0.982 \end{pmatrix}$$

- Output: $f : V \rightarrow \mathbb{R}^d$


Word2vec

word = "sweden"


Word	Cosine distance

norway	0.760124
denmark	0.715460
finland	0.620022
switzerland	0.588132
belgium	0.585835
netherlands	0.574631
iceland	0.562368
estonia	0.547621
slovenia	0.531408

Word2vec




Continuous Bag of Words (CBOW)




Skip-grams

Skip-gram

- The idea: we want to use words to **predict** their context words
- Context: a fixed window of size $2m$



Skip-gram



Skip-gram: objective function

- For each position $t = 1, 2, \dots, T$, predict context words within context size m , given center word w_j :

$$\mathcal{L}(\theta) = \prod_{t=1}^T \prod_{-m \leq j \leq m, j \neq 0} P(w_{t+j} \mid w_t; \theta)$$

all the parameters to be optimized



- The objective function $J(\theta)$ is the (average) negative log likelihood:

$$J(\theta) = -\frac{1}{T} \log \mathcal{L}(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} \mid w_t; \theta)$$

How to define $P(w_{t+j} \mid w_t; \theta)$?

- We have two sets of vectors for each word in the vocabulary

$\mathbf{u}_i \in \mathbb{R}^d$: embedding for target word i

$\mathbf{v}_{i'} \in \mathbb{R}^d$: embedding for context word i'

- Use inner product $\mathbf{u}_i \cdot \mathbf{v}_{i'}$ to measure how likely word i appears with context word i' , the larger the better

“softmax” we learned last time!

$$P(w_{t+j} \mid w_t) = \frac{\exp(\mathbf{u}_{w_t} \cdot \mathbf{v}_{w_{t+j}})}{\sum_{k \in V} \exp(\mathbf{u}_{w_t} \cdot \mathbf{v}_k)}$$


$\theta = \{\{\mathbf{u}_k\}, \{\mathbf{v}_k\}\}$ are all the parameters in this model!

How to train the model

Calculating all the gradients together!

$$\theta = \{\{\mathbf{u}_k\}, \{\mathbf{v}_k\}\}$$

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} \mid w_t; \theta) \quad \nabla_{\theta} J(\theta) = ?$$

Q: How many parameters are in total?

We can apply stochastic gradient descent (SGD)!

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla_{\theta} J(\theta)$$

Skip-gram with negative sampling (SGNS)

Idea: recast problem as binary classification!

- Target word is positive example
- All words not in context are negative

$$P(D = 1 \mid t, c) = \sigma(\mathbf{u}_t \cdot \mathbf{v}_c)$$


$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$

positive examples +

t	c
apricot	tablespoon
apricot	of
apricot	jam
apricot	a

negative examples -


t	c	t	c
apricot	aardvark	apricot	seven
apricot	my	apricot	forever
apricot	where	apricot	dear
apricot	coaxial	apricot	if



To compute loss, pick K random words as negative examples:

$$J(\theta) = -P(D = 1 \mid t, c) - \frac{1}{K} \sum_{i=1}^K P(D = 0 \mid t_i, c)$$

Continuous Bag of Words (CBOW)



$$L(\theta) = \prod_{t=1}^T P(w_t \mid \{w_{t+j}\}, -m \leq j \leq m, j \neq 0)$$

$$\bar{\mathbf{v}}_t = \frac{1}{2m} \sum_{-m \leq j \leq m, j \neq 0} \mathbf{v}_{t+j}$$


$$P(w_t \mid \{w_{t+j}\}) = \frac{\exp(\mathbf{u}_{w_t} \cdot \bar{\mathbf{v}}_t)}{\sum_{k \in V} \exp(\mathbf{u}_k \cdot \bar{\mathbf{v}}_t)}$$

GloVe: **G**lobal **V**ectors

- Let's take the global co-occurrence statistics: $X_{i,j}$

$$J = \sum_{i,j=1}^V f(X_{ij}) \left(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij} \right)^2$$

- Training faster
- Scalable to very large corpora



(Pennington et al, 2014): GloVe: Global Vectors for Word Representation

GloVe: **G**lobal **V**ectors

Nearest words to
frog:

1. frogs
2. toad
3. litoria
4. leptodactylidae
5. rana
6. lizard
7. eleutherodactylus



litoria



leptodactylidae



rana



eleutherodactylus

(Pennington et al, 2014): GloVe: Global Vectors for
Word Representation

FastText: Sub-Word Embeddings

- Similar as Skip-gram, but break words into n-grams with $n = 3$ to 6

where: 3-grams: <wh, whe, her, ere, re>

4-grams: <whe, wher, here, ere>

5-grams: <wher, where, here>

6-grams: <where, where>

- Replace $\mathbf{u}_i \cdot \mathbf{v}_j$ by $\sum_{g \in n\text{-grams}(w_i)} \mathbf{u}_g \cdot \mathbf{v}_j$

- More to come! Contextualized word embeddings



FastText: Sub-Word Embeddings

- Similar as Skip-gram, but break words into n-grams with $n = 3$ to 6

where: 3-grams: <wh, whe, her, ere, re>

4-grams: <whe, wher, here, ere>

5-grams: <wher, where, here>

6-grams: <where, where>

- Replace $\mathbf{u}_i \cdot \mathbf{v}_j$ by $\sum_{g \in n\text{-grams}(w_i)} \mathbf{u}_g \cdot \mathbf{v}_j$

- More to come! Contextualized word embeddings



(Bojanowski et al, 2017): Enriching Word Vectors with Subword Information

Trained word embeddings available

- word2vec: <https://code.google.com/archive/p/word2vec/>
- GloVe: <https://nlp.stanford.edu/projects/glove/>
- FastText: <https://fasttext.cc/>

Download pre-trained word vectors

- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License](http://www.opendatacommons.org/licenses/pddl/1.0/) v1.0 whose full text can be found at: <http://www.opendatacommons.org/licenses/pddl/1.0/>.
 - [Wikipedia 2014 + Gigaword 5](#) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download): [glove.6B.zip](#)
 - Common Crawl (42B tokens, 1.9M vocab, uncased, 300d vectors, 1.75 GB download): [glove.42B.300d.zip](#)
 - Common Crawl (840B tokens, 2.2M vocab, cased, 300d vectors, 2.03 GB download): [glove.840B.300d.zip](#)
 - Twitter (2B tweets, 27B tokens, 1.2M vocab, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download): [glove.twitter.27B.zip](#)
- Ruby [script](#) for preprocessing Twitter data


Differ in algorithms, text corpora, dimensions, cased/uncased...

Evaluating Word Embeddings

Extrinsic vs intrinsic evaluation

Extrinsic evaluation

- Let's plug these word embeddings into a real NLP system and see whether this improves performance
- Could take a long time but still the most important evaluation metric



Intrinsic evaluation

- Evaluate on a specific/intermediate subtask
- Fast to compute
- Not clear if it really helps the downstream task

Intrinsic evaluation

Word similarity

Example dataset: wordsim-353

353 pairs of words with human judgement

<http://www.cs.technion.ac.il/~gabr/resources/data/wordsim353/>

Word 1	Word 2	Human (mean)
tiger	cat	7.35
tiger	tiger	10
book	paper	7.46
computer	internet	7.58
plane	car	5.77
professor	doctor	6.62
stock	phone	1.62
stock	CD	1.31
stock	jaguar	0.92

Cosine similarity:

$$\cos(\mathbf{u}_i, \mathbf{u}_j) = \frac{\mathbf{u}_i \cdot \mathbf{u}_j}{\|\mathbf{u}_i\|_2 \times \|\mathbf{u}_j\|_2}.$$

Metric: Spearman rank correlation

Intrinsic evaluation

Word Similarity

Model	Size	WS353	MC	RG	SCWS	RW
SVD	6B	35.3	35.1	42.5	38.3	25.6
SVD-S	6B	56.5	71.5	71.0	53.6	34.7
SVD-L	6B	65.7	<u>72.7</u>	75.1	56.5	37.0
CBOW [†]	6B	57.2	65.6	68.2	57.0	32.5
SG [†]	6B	62.8	65.2	69.7	<u>58.1</u>	37.2
GloVe	6B	<u>65.8</u>	<u>72.7</u>	<u>77.8</u>	53.9	<u>38.1</u>
SVD-L	42B	74.0	76.4	74.1	58.3	39.9
GloVe	42B	<u>75.9</u>	<u>83.6</u>	<u>82.9</u>	<u>59.6</u>	<u>47.8</u>
CBOW*	100B	68.4	79.6	75.4	59.4	45.5

Intrinsic evaluation

Word analogy

man: woman $\approx$ king: ?

$$\arg \max_i (\cos(\mathbf{u}_i, \mathbf{u}_b - \mathbf{u}_a + \mathbf{u}_c))$$

semantic

syntactic

Chicago:Illinois $\approx$ Philadelphia: ? bad:worst $\approx$ cool: ?

More examples at

<http://download.tensorflow.org/data/questions-words.txt>

What can go wrong with word embeddings?

- What's wrong with learning a word's "meaning" from its usage?
- What data are we learning from?
- What are we going to learn from this data?

What do we mean by bias?

- Identify *she* - *he* axis in word vector space, project words onto this axis
- Nearest neighbor of $(b - a + c)$

Extreme *she* occupations

- | | | |
|-----------------|-----------------------|------------------------|
| 1. homemaker | 2. nurse | 3. receptionist |
| 4. librarian | 5. socialite | 6. hairdresser |
| 7. nanny | 8. bookkeeper | 9. stylist |
| 10. housekeeper | 11. interior designer | 12. guidance counselor |

Extreme *he* occupations

- | | | |
|----------------|-------------------|----------------|
| 1. maestro | 2. skipper | 3. protege |
| 4. philosopher | 5. captain | 6. architect |
| 7. financier | 8. warrior | 9. broadcaster |
| 10. magician | 11. fighter pilot | 12. boss |


Bolukbasi et al. (2016)

Racial Analogies	
black → homeless	caucasian → servicemen
caucasian → hillbilly	asian → suburban
asian → laborer	black → landowner
Religious Analogies	
jew → greedy	muslim → powerless
christian → familial	muslim → warzone
muslim → uneducated	christian → intellectually

Manzini et al. (2019)

Debiasing


- Identify gender subspace with gendered words
- Project words onto this subspace
- Subtract those projections from the original word



Bolukbasi et al. (2016)


Word embedding applications

- Can learn to map multiple kinds of data into a single representation.
- E.g., bilingual English and Mandarin Chinese word-embedding as in [Socher et al. \(2013a\)](#).
- Embed as above, but words that are known as close translations should be close together.
- Words we *didn't know* were translations end up close together!
- Structures of two languages get pulled into alignment.




Word embedding applications

- Can apply to get a joint embedding of words and images or other multi-modal data sets.
- New classes map near similar existing classes: e.g., if 'cat' is unknown, cat images map near dog.



(Socher *et al.* (2013b))

<http://colah.github.io/posts/2014-07-NLP-RNNs-Representations/>



Language Modeling: predict the next word

Assign probabilities to text.

Given a sequence $(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T)$, we want to **maximize** $P(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T)$.

$$P(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T) = P(\mathbf{x}_1)P(\mathbf{x}_2|\mathbf{x}_1)P(\mathbf{x}_3|\mathbf{x}_1, \mathbf{x}_2)P(\mathbf{x}_4|\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3) \dots P(\mathbf{x}_T|\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{T-1})$$

$P(\text{I like cats because they look cute}) = P(\text{I}) P(\text{like} | \text{I}) P(\text{cats} | \text{I like}) P(\text{as} | \text{I like cats}) P(\text{they} | \text{I like cats because})$

$P(\text{look} | \text{I like cats because they}) P(\text{cute} | \text{I like cats because they look})$

Predict the next word given current text!

n -Gram Language Model

n -Gram: chunk of n consecutive words

Count the frequency of each n -grams and predict next word!

Uni-gram: "I" "like" "cats" "as" "they" "look" "cute"

Bi-gram: "I like" "like cats" "cats as" "as they" ...

Tri-gram: "I like cats" "like cats as" "cats as they" ...

Assume each word only depends on previous $n - 1$ words.

$$\begin{aligned} P(\mathbf{x}_t | \mathbf{x}_1, \dots, \mathbf{x}_{t-1}) &= P(\mathbf{x}_t | \mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1}) \\ &= \frac{\text{count}(\mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1}, \mathbf{x}_t)}{\text{count}(\mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1})} \end{aligned}$$

In *bi-gram* LM

$P(\text{I like cats as they look cute}) = P(\text{I}) P(\text{like} | \text{I}) P(\text{cats} | \text{like}) P(\text{as} | \text{cats}) P(\text{they} | \text{because}) P(\text{look} | \text{they}) P(\text{cute} | \text{look})$

n -Gram Language Model

n -Gram: chunk of n consecutive words

Count the frequency of each n -grams and predict next word!

Uni-gram: “I” “like” “cats” “as” “they” “look” “cute”

Bi-gram: “I like” “like cats” “cats as” “as they” ...

Tri-gram: “I like cats” “like cats as” “cats as they” ...

Assume each word only depends on previous $n - 1$ words.

$$\begin{aligned} P(\mathbf{x}_t | \mathbf{x}_1, \dots, \mathbf{x}_{t-1}) &= P(\mathbf{x}_t | \mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1}) \\ &= \frac{\text{count}(\mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1}, \mathbf{x}_t)}{\text{count}(\mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1})} \end{aligned}$$

In *bi-gram* LM

$P(\text{I like cats as they look cute}) = P(\text{I}) P(\text{like} | \text{I}) P(\text{cats} | \text{like}) P(\text{as} | \text{cats}) P(\text{they} | \text{because}) P(\text{look} | \text{they}) P(\text{cute} | \text{look})$

n -Gram Language Model: issue

n -Gram: chunk of n consecutive words

Uni-gram: "I" "like" "cats" "as" "they" "look" "cute"

Bi-gram: "I like" "like cats" "cats as" "as they" ...

Tri-gram: "I like cats" "like cats as" "cats as they" ...

Count the frequency of each n -grams and predict next word!

Assume each word only depends on previous $n - 1$ words.

$$\begin{aligned} P(\mathbf{x}_t | \mathbf{x}_1, \dots, \mathbf{x}_{t-1}) &= P(\mathbf{x}_t | \mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1}) \\ &= \frac{\text{count}(\mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1}, \mathbf{x}_t)}{\text{count}(\mathbf{x}_{t-n+1}, \dots, \mathbf{x}_{t-1})} \end{aligned}$$

Increasing n provides contextual information, but exponentially increases the size of the counting table!

Bag of Words (to represent documents)

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



it	6
I	5
the	4
to	3
and	3
seen	2
yet	1
would	1
whimsical	1
times	1
sweet	1
satirical	1
adventure	1
genre	1
fairy	1
humor	1
have	1
great	1
...	...

Drawbacks:

- High dimensionality
- No semantic information

Document similarity?

document 1

Obama
speaks
to
the
media
in
Illinois



Documents have no words in common.
How can we quantify that they convey
similar meanings?
(Assume B. Obama is president.)

document 2

The
President
greets
the
press
in
Chicago


Semantic similarity

- Motivation
 - Put words into vectors so we can measure the similarity between words
 - Use cosine similarity



Why Do We Need Word Embeddings?

- Why Do We Need Word Embeddings?
 - Numerical Input
 - Shows Similarity and Distance




	living being	feline	human	gender	royalty	verb	plural
cat	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
kitten	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
dog	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
houses	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8
man	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
woman	0.7	0.3	0.9	-0.7	0.1	-0.5	-0.4
king	0.5	-0.4	0.7	0.8	0.9	-0.7	-0.6
queen	0.8	-0.1	0.8	-0.9	0.8	-0.5	-0.9

embedding using features of words

What are word embeddings

- What are Word Embeddings?
 - vector representations of words that capture semantic relationships
 - Latent Semantic Analysis / Indexing [S. Deerwester et al 1988]



Efficient Estimation of Word Representations in Vector Space

Tomas Mikolov

Google Inc., Mountain View, CA
tmikolov@google.com

Kai Chen

Google Inc., Mountain View, CA
kaichen@google.com

Greg Corrado

Google Inc., Mountain View, CA
gcorrado@google.com

Jeffrey Dean

Google Inc., Mountain View, CA
jeff@google.com

Abstract

We propose two novel model architectures for computing continuous vector representations of words from very large data sets. The quality of these representations is measured in a word similarity task, and the results are compared to the previously best performing techniques based on different types of neural networks. We observe large improvements in accuracy at much lower computational cost, i.e. it takes less than a day to learn high quality word vectors from a 1.6 billion words data set. Furthermore, we show that these vectors provide state-of-the-art performance on our test set for measuring syntactic and semantic word similarities.



X 2 vec

- Generate vector representations (embeddings) for various data types
- Examples:
 - Word2Vec
 - Doc2Vec
 - Node2Vec
 - Item2Vec
 - Sent2Vec



Doc2Vec

- A vector to represent a paragraph, regardless of length
 - embeddings for paragraph and words
 - Applications: Document classification, sentiment analysis, recommendation systems, and information retrieval




Problems with word2vec

- Words with multiple meanings only have one representation
 - eg. **bank** of river or **bank** of money
 - Need contextual information
- Limited Context
 - only trained on words within the context window




How to use word vectors with neural networks?




- Inputs and outputs don't have fixed lengths
- Features are not shared

Let's simplify!


What if we have a single word and a single output?




Recurrent neural network (RNN)




Recurrent neural network (RNN)



Recurrent neural network (RNN)



Parameterize RNN



- Too many parameters if we have a long sequence!
- Longer sequence parameters will not receive many updates
- What if sequence lengths vary?

RNN w/ parameter-sharing

Simple fix: use **the same parameters** across different timesteps.

